

ESE 507

Project: Part 1

Jin Yuan Chen
115702978

October 8, 2026

Contents

1	Multiply Accumulate (MAC) Design	1
1.1	Design Parameter	1
2	Combinational Synthesis Results	2
2.1	Area	2
2.2	Power	3
2.3	Critical Paths	3
3	Pipelined Synthesis Results	4
3.1	Area	4
3.2	Power	5
3.3	Critical Paths	6
4	Performance Comparison	7
4.1	Maximum-Frequency Energy	7
4.2	Frequency-Dependent Energy	7
4.3	Latency and Throughput	7
4.4	Design Trade-offs	7
5	Potential for Deeper Pipelining	7
5.1	Multiplier Pipelining	7
5.2	Adder Pipelining	7
6	Variable Parameter Synthesis Results	7
6.1	Accumulator Width (ACCW)	7
6.2	Numerical-Bit Width (WIDTH)	7
6.3	Scaling Behavior	7
6.4	Accumulator Overflow	7
7	Appendix	7

1. Multiply Accumulate (MAC) Design

Two versions of the MAC were implemented to compare the unpipelined and pipelined designs. Both the unpipelined `mac` and pipelined `mac_pipe` designs use an accumulator register that updates its output on the rising edge of the clock. The primary difference between the two designs is the combinational multiply-add operation. In the unpipelined `mac`, the multiplication and addition are performed within a single combinational stage before the result is stored in the accumulator. In the pipelined `mac_pipe`, the multiply-add operation is divided into two stages. The multiplier output and associated control signals are stored in an additional intermediate register before the addition is performed and the result is stored in the accumulator on the following clock cycle. Consequently, the pipelined design introduces an additional cycle of latency, requiring two clock cycles for an input to propagate to the accumulator output, compared with one cycle for the unpipelined design. However, the benefits of pipelining become apparent when considering the overall throughput and frequency of the system.

1.1. Design Parameter

Both the `mac` and `mac_pipe` designs were synthesized with `WIDTH = 16` and `ACCW = 48` across a range of clock frequencies. The analysis spans frequencies from a relatively low frequency of 10 MHz to a maximum of 1 GHz. At each frequency, the area, power, critical path location, and timing constraint of each design are evaluated.


2. Combinational Synthesis Results

The unpipelined `mac` design was synthesized to determine the maximum achievable clock frequency with `WIDTH = 16` and `ACCW = 48`. The synthesis workflow was automated with scripts that write the configured design parameters and the synthesis script, before running synthesis on the remote CAD server. The generated synthesis report is filtered for key attributes, including area, power, and timing. For convenient data analysis, these attributes are stored in comma-separated value (CSV) files for every synthesis run. The tabular structure of the CSV files is well suited for Python plotting libraries, which were used to generate the graphs below. For more details on the automation workflow, refer to `scr/`. Only data points that **meet the timing constraint** are plotted in the graphs below.

2.1. Area

The design area is estimated from the synthesis report as the total cell area, which does not account for net interconnect area. Therefore, the actual physical area of the design with interconnects may be larger than the reported area. Further analysis of design area uses the conservative estimation as the sum of combinational area, buffer/inverter area, noncombinational area and macro/black box area in the synthesis report.

As shown in Figure 1, the total design area remains relatively constant as the frequency increases. However, an overall upward trend is seen, particularly at frequencies above 200 MHz. This increase can be attributed to the additional area required to meet the more difficult timing constraints associated with higher frequency. The approximately linear increase in area beyond 200 MHz indicates the achievable trade-off for frequency at the expensive of area. In contrast, the relatively constant area at lower frequencies suggests that the synthesis tool implemented the design to a near-minimal cell configuration regardless of less restrictive timing constraints.



Freq. (MHz)	Area (μm^2)
10	2067.09
20	2067.09
25	2067.09
50	2067.09
100	2067.09
166.67	2067.09
200	2066.29
333.33	2164.18
:	:
500	2320.85
526.32	2300.90
555.56	2302.76
769.23	2424.06


Figure 1: Area, frequency for unpipelined `mac` (MET).

Table 1: Selected MET area points for `mac`.

An interesting pattern is the occasional small downward trend in area at higher frequencies. This suggests that, at certain frequency, the synthesis tool may identify a simplified version that meets timing constraint with less cell area. For example, the design synthesized at 555.56 MHz has a slightly lower total area than the design synthesized at 500 MHz, despite the overall increasing trend of frequency and area. This behavior indicates that the relationship between frequency and synthesized area is not strictly monotonic and can depend on the specific optimizations selected by the synthesis tool.

2.2. Power

The Total Power shown in Figure 2 is defined as the sum of Total Dynamic Power and Cell Leakage Power in the synthesis report. The Total Dynamic Power is also reported as the sum of Cell Internal Power and Net Switching Power, both contributing an near equal portion of the sum.



Freq. (MHz)	Σ Power (μW)
10	65.49
20	83.85
25	93.03
50	138.92
100	230.72
166.67	353.12
200	411.86
333.33	698.23
...	...
500	1127.93
526.32	1148.45
555.56	1218.61
769.23	1745.19

Figure 2: Total power, frequency for unpipelined mac (MET).

Table 2: Selected MET power points for mac.

The Cell Leakage Power is almost negligible compared with the dynamic power. This may be attributed to the relatively small area of the design, with the accumulator register being the primary memory elements contributing to leakage power. In the case that the frequency is considerably slow, the dynamic power can become sufficiently small that leakage power becomes a significant contributor to total power. Otherwise, given a fixed accumulator width and integer-bit width, the leakage power remains approximately constant at frequency below 166.67 MHz, as shown in the table of Figure 1. This suggests that leakage power is primarily dependent on the area of the design. The slight increase in leakage power at higher frequencies can therefore be attributed to the corresponding increase in design area.

Consequently, the trend seen in Figure 2 for total power is largely influenced by the Dynamic Power and its dependence on frequency. Increasing the frequency results in a near linear increase in Dynamic Power, with an approximate power increase of $1.836 \mu W/MHz$. In other words, a higher-frequency design performs more switching operation per unit time than a lower-frequency design, resulting in higher power consumption over the same period of time.

2.3. Critical Paths

The synthesis report provides the start and end points of the critical paths, along with the paths themselves and their data arrival times. If the data arrival time is less than or equal to the data required time, the design constraint is considered MET, as indicated by Slack (MET) in the report.

A design with Slack (VIOLATED) indicates a timing constraint violation and will typically display a negative slack value, calculated as the difference between the data arrival time and the data required time. A negative value indicates that the timing constraint has not been met. A design reported as Slack (VIOLATED: increase significant digits) is also considered a violation, even if the displayed slack value is zero. This occurs because the actual slack value is slightly negative but is rounded to zero due to the limited number of significant digits displayed in the report.

Freq. (MHz)	Area (μm^2)	Power (μW)	Slack (ns)	Path Start	Path End	Timing
10	2067.09	65.49	95.68	input0[1]	acc_reg[47]	MET
11.11	2067.09	67.53	85.68	input0[1]	acc_reg[47]	MET
12.5	2067.09	70.08	75.68	input0[1]	acc_reg[47]	MET
14.29	2067.09	73.35	65.68	input0[1]	acc_reg[47]	MET
16.67	2067.09	77.73	55.68	input0[1]	acc_reg[47]	MET
20	2067.09	83.85	45.68	input0[1]	acc_reg[47]	MET
25	2067.09	93.03	35.68	input0[1]	acc_reg[47]	MET
33.33	2067.09	108.32	25.68	input0[1]	acc_reg[47]	MET
50	2067.09	138.92	15.68	input0[1]	acc_reg[47]	MET
100	2067.09	230.72	5.68	input0[1]	acc_reg[47]	MET
111.11	2067.09	251.12	4.68	input0[1]	acc_reg[47]	MET
125	2067.09	276.62	3.68	input0[1]	acc_reg[47]	MET
142.86	2067.09	309.40	2.68	input0[1]	acc_reg[47]	MET
166.67	2067.09	353.12	1.68	input0[1]	acc_reg[47]	MET
200	2066.29	411.86	0.68	input0[1]	acc_reg[47]	MET
250	2102.20	525.46	0.01	input0[0]	acc_reg[47]	MET
333.33	2164.18	698.23	0.00	input0[7]	acc_reg[47]	MET
500	2320.85	1127.93	0.00	input0[3]	acc_reg[47]	MET
526.32	2300.90	1148.45	0.03	input0[5]	acc_reg[47]	MET
555.56	2302.76	1218.61	0.00	input0[7]	acc_reg[47]	MET
588.24	2295.58	1277.55	0.01	input0[11]	acc_reg[47]	MET
625	2321.12	1369.41	0.00	input0[5]	acc_reg[47]	MET
666.67	2341.60	1485.30	0.00	input0[5]	acc_reg[44]	MET
714.29	2395.60	1626.06	0.00	input0[0]	acc_reg[35]	MET
769.23	2424.06	1745.19	0.00	input0[7]	acc_reg[31]	MET
833.33	2471.14	2003.27	-0.07	input0[5]	acc_reg[33]	VIOLATED
909.09	2543.23	2229.30	-0.14	input0[15]	acc_reg[21]	VIOLATED
1000	2529.13	2441.94	-0.22	input0[7]	acc_reg[38]	VIOLATED

Table 3: Frequency, area, power, timing slack, and critical path for unpipelined mac.

The critical path starts from an `input0` register and ends at `acc_reg`, with slight variations in the indices of `input0` and `acc_reg`. At low frequencies and up to approximately 200 MHz, the critical path remains the same, starting from `input0[1]` and ending at `acc_reg[47]`. Between 200 MHz and 625 MHz, the endpoint of the critical path remains fixed at `acc_reg[47]`, while the index of `input0` varies depending on the frequency. Above 625 MHz, both the `input0` index and the endpoint index of `acc_reg` vary without a clear or consistent pattern as shown in Table 3.

In particular, there appears to be a tendency for the critical path to shift toward higher-indexed `input0` elements and, in other cases, lower-indexed `acc_reg` elements. This indicate a possible trend in which the synthesis fail to optimize the design of higher indexed paths, in contrast to the critical path at lower frequency to be at lower indexed of `input0`. Another explanation is that the paths between different indexed elements have the same critical path or delay, which synthesis does not differentiate. Therefore, additional analysis is required to determine whether this trend is directly caused by the stricter timing constraints, the design area, or their interaction.

3. Pipelined Synthesis Results

In a similar fashion, the pipelined `mac_pipe` design was synthesized to determine the maximum achievable clock frequency with `WIDTH = 16` and `ACCW = 48`. The same automated scripts were used to collect and store the synthesis results in a CSV file. The resulting datasets for the unpipelined `mac` and pipelined `mac_pipe` designs were then analyzed to compare their area, power, and timing performance. In terms of performance metric, the maximum achievable frequency was prioritized before considering tradeoff in area and power, whose relative importance depends on the target application. Only data points that meet the timing constraint are plotted in the graphs below.

3.1. Area

The design area of the pipelined `mac_pipe` is the estimated cell area from the synthesis report. For better comparison, the area for both the unpipelined and pipelined designs are plotted over frequency.




Figure 3: Area between mac and mac_pipe (MET).

Freq. (MHz)	Area (μm^2)		ΔArea
	mac	mac_pipe	
10	2067.09	2306.75	239.67
20	2067.09	2306.49	239.40
25	2067.09	2306.49	239.40
50	2067.09	2306.49	239.40
100	2067.09	2306.49	239.40
166.67	2067.09	2306.49	239.40
200	2066.29	2305.16	238.87
333.33	2164.18	2459.44	295.26
⋮	⋮	⋮	⋮
500	2320.85	2648.03	327.18
526.32	2300.90	2639.78	338.88
555.56	2302.76	2692.98	390.22
769.23	2424.06	2773.58	349.52

Table 4: Selected MET area points.

As shown in Figure 3, the area of the pipelined design is generally higher than the unpipelined design, which is expected given the additional pipeline registers. Accounting for increasing in frequency, the area of the pipelined design is about $240 \mu\text{m}^2$ higher than the unpipelined design at 10 MHz, increasing to about $320 \mu\text{m}^2$ at 500 MHz. This trend suggests that the pipelined design is area-deficient at higher frequencies.

An interesting trend in area occurs at the particular frequency from 500MHz to 700MHz. At those frequency the area of the pipelined design increase while the unpipelined design decrease, such that the area difference is significantly greater than at other frequencies. This suggests that pipelined design is more area-deficient at frequencies where the unpipelined design could be optimized to be more area-efficient.

3.2. Power

The Total Power is defined the same way as the unpipelined design. Previously, dynamic power was analyzed to cause the trend seen in Figure 2. This trend is also seen in Figure 4 for the pipelined design.




Figure 4: Power between mac and mac_pipe (MET).

Freq. (MHz)	$\sum$ Power (μW)		ΔPower
	mac	mac_pipe	
10	65.49	71.03	5.55
20	83.85	90.75	6.90
25	93.03	100.61	7.59
50	138.92	149.92	10.99
100	230.72	248.53	17.81
166.67	353.12	380.01	26.89
200	411.86	446.49	34.63
333.33	698.23	744.59	46.36
⋮	⋮	⋮	⋮
500	1127.93	1182.18	54.26
526.32	1148.45	1239.58	91.13
555.56	1218.61	1333.95	115.34
769.23	1745.19	1928.30	183.11

Table 5: Selected MET power points.

Cell Leakage Power continues to be almost negligible compared with dynamic power, as in the unpipelined design. The slight increase in leakage power for the pipelined design is attributed to the additional pipeline registers increasing the cell area over which leakage can occur. Thus the leakage power could be a more significant fraction in larger designs, especially at lower frequencies. Otherwise, leakage stays approximately constant at low frequency as seen in unpipelined design with the same parameters. The slight increase in power at higher frequency can be attributed to the increased design area from synthesis.

The power difference is minimal at lower frequency, but growing at a larger rate than the unpipelined design. This can be a consequence of the larger pipelined area providing more cells that will be switched each cycle. The rate of power growth becomes more distinct at higher frequency. For example, Δ Power remains on the order of tens of microwatts through 500 MHz and then continues to peak, reaching about 183 μ W at 769.23 MHz. At this rate, the pipelined design may significantly exceed the power budget compared to the unpipelined design at increasing frequency.

3.3. Critical Paths

The start and end points of the critical path for `mac_pipe` are reported in Table 6.

Because the multiply-add operation is split across two registers, the critical path may lie in either the multiplier stage (`reg1`) or the accumulator stage (`reg2`).

Freq. (MHz)	Area (μ m ²)	Power (μ W)	Slack (ns)	Path Start	Path End	Timing
10	2306.75	71.03	95.74	reg2/out_accum_reg[0]	reg2/out_accum_reg[47]	MET
11.11	2306.75	73.22	85.74	reg2/out_accum_reg[0]	reg2/out_accum_reg[47]	MET
12.5	2306.49	75.96	75.74	reg2/out_accum_reg[0]	reg2/out_accum_reg[47]	MET
14.29	2306.75	79.48	65.74	reg2/out_accum_reg[0]	reg2/out_accum_reg[47]	MET
16.67	2306.75	84.18	55.74	reg2/out_accum_reg[0]	reg2/out_accum_reg[47]	MET
20	2306.49	90.75	45.74	reg2/out_accum_reg[0]	reg2/out_accum_reg[47]	MET
25	2306.49	100.61	35.74	reg2/out_accum_reg[0]	reg2/out_accum_reg[47]	MET
33.33	2306.49	117.05	25.74	reg2/out_accum_reg[0]	reg2/out_accum_reg[47]	MET
50	2306.49	149.92	15.74	reg2/out_accum_reg[0]	reg2/out_accum_reg[47]	MET
100	2306.49	248.53	5.74	reg2/out_accum_reg[0]	reg2/out_accum_reg[47]	MET
111.11	2306.49	270.44	4.74	reg2/out_accum_reg[0]	reg2/out_accum_reg[47]	MET
125	2306.49	297.83	3.74	reg2/out_accum_reg[0]	reg2/out_accum_reg[47]	MET
142.86	2306.49	333.05	2.74	reg2/out_accum_reg[0]	reg2/out_accum_reg[47]	MET
166.67	2306.49	380.01	1.74	reg2/out_accum_reg[0]	reg2/out_accum_reg[47]	MET
200	2305.16	446.49	0.75	reg2/out_accum_reg[0]	reg2/out_accum_reg[47]	MET
250	2342.93	554.37	0.12	reg1/out_product_reg[0]	reg2/out_accum_reg[47]	MET
333.33	2459.44	744.59	0.00	reg2/out_accum_reg[0]	reg2/out_accum_reg[47]	MET
500	2648.03	1182.18	0.00	reg2/out_accum_reg[9]	reg2/out_accum_reg[47]	MET
526.32	2639.78	1239.58	0.00	reg2/out_accum_reg[7]	reg2/out_accum_reg[46]	MET
555.56	2692.98	1333.95	0.00	input0[7]	reg1/out_product_reg[30]	MET
588.24	2693.52	1407.20	0.01	input0[3]	reg1/out_product_reg[30]	MET
625	2722.78	1541.62	0.00	input0[5]	reg1/out_product_reg[30]	MET
666.67	2732.35	1636.92	0.00	input0[5]	reg1/out_product_reg[29]	MET
714.29	2724.64	1741.75	0.00	input0[3]	reg1/out_product_reg[30]	MET
769.23	2773.58	1928.30	0.00	reg2/out_accum_reg[7]	reg2/out_accum_reg[47]	MET
833.33	2823.86	2116.52	0.00	input0[11]	reg1/out_product_reg[23]	VIOLATED*
909.09	2860.56	2372.62	-0.06	input0[13]	reg1/out_product_reg[27]	VIOLATED
1000	2832.37	2581.64	-0.18	input0[9]	reg1/out_product_reg[20]	VIOLATED

Table 6: Frequency, area, power, timing slack, and critical path for pipelined `mac_pipe`.

At low frequencies and up to approximately 200 MHz, the critical path is consistently from bit [0] to bit [47] of `out_accum_reg` in `reg2`. This is the 48-bit adder in the second pipeline stage: the carry chain from the least-significant bit of the accumulator to the most-significant bit. At 250 MHz the start point shifts to bit [0] of `out_product_reg` in `reg1`, while the endpoint remains bit [47] of `out_accum_reg` in `reg2`. The path is still in the accumulate stage, now starting from the registered product rather than an accumulator bit. From 333.33 MHz through 526.32 MHz the path remains in the accumulate stage, with some variation in the start and end indices of `out_accum_reg`.

Between 555.56 MHz and 714.29 MHz the critical path moves to the multiplier stage, starting from an `input0` bit and ending at `reg1/out_product_reg`. At 769.23 MHz, the last frequency that meets timing, the path returns to the accumulate stage. Timing violations from 833.33 MHz onward are again on the

multiplier path, from `input0` to `reg1/out_product_reg`. The 833.33 MHz point has a displayed slack of zero but is reported as VIOLATED, which is the same rounded-slack case described for the unpipelined design.

This behavior follows from how the design is pipelined. The combinational multiply `input0 * input1` is captured in `reg1`, and the combinational add `out_product + post_accum` is captured in `reg2`. At relaxed periods the longest delay is the accumulator adder, so the tool reports a path through `reg2`. As the period shrinks, the 16×16 multiplier delay becomes comparable to, and then larger than, the adder delay, so the critical path shifts into `reg1`. The brief return to the accumulate stage at 769.23 MHz indicates that the two stages are similarly tight near the timing limit. That the first violations occur on the multiplier path suggests that the product stage, rather than the 48-bit adder, is what ultimately limits the maximum frequency of `mac_pipe`.

4. Performance Comparison

4.1. Maximum-Frequency Energy

4.2. Frequency-Dependent Energy

4.3. Latency and Throughput

4.4. Design Trade-offs

5. Potential for Deeper Pipelining

5.1. Multiplier Pipelining

5.2. Adder Pipelining

6. Variable Parameter Synthesis Results

6.1. Accumulator Width (ACCW)

6.2. Numerical-Bit Width (WIDTH)

6.3. Scaling Behavior

6.4. Accumulator Overflow

7. Appendix